


# Database Administration

- Database Administration
- Your PostgreSQL DBMSs
- Catalogs
- Catalogs
- Access to System Catalog
- PostgreSQL Catalog
- Security, Privilege, Authorisation
- Database Security
- Database Access Control
- Database Access Control in PostgreSQL
- SQL Access Control
- SQL Access Control in PostgreSQL
- Problems with SQL Access Control
- Mandatory Access Control
- Security in Statistical Databases
- Performance Tuning
- Performance Tuning
- Denormalisation
- Indexes
- Query Tuning
- PostgreSQL Query Tuning
- EXPLAIN Examples

# Database Administration

Database installations are typically ...

- used by **many** casual users (typically via a web interface)
- used by **few** developers (who build schemas and applications)
- managed by **one** DB administrator (or a small team of DBAs)



Managing DBMS installation

Building schemas/queries/...

Accessing DB using parametric queries via e.g. web interface

## Database Administration (cont)

Tasks of the database administrator:

- manage the database installation  
(including configuring system parameters, running the server, back-ups, ...)
  - create users and specify what they can/cannot do  
(according to the security policies of the organisation running the DBMS)
  - tune performance of applications and overall system
  - act as the owner of the system meta-data (e.g. catalog)
  - back-up the contents of the database (in case of disasters)
- 

## Database Administration (cont)

The DBMS itself assists DB administration by:

- maintaining a database of schemas in the system (catalog)
  - maintaining a database of users/groups and their privileges  
(which determines who can perform which operations on which objects)
  - maintaining statistical information about database instances  
(used by the query optimiser in determining best query execution strategy)
-

# Your PostgreSQL DBMSs

In the PostgreSQL installations on **grieg** you are the DBA  
(aka PostgreSQL super-user, owner of postmaster process)

Data/files associated with server live under **/srvr/YOU/pgsql1903/**

Some things that you could potentially do:

- change configuration parameters of your sever  
(e.g. more run-time buffers, more connections, **postgresql.conf**)
- add new users of the database and set their privileges  
(SQL commands: **CREATE USER**, **CREATE ROLE**, **GRANT**, **REVOKE**)
- modify accessibility of the databases  
(**pg\_hba.conf**, who can access what from where and authentication)

---

## § Catalogs

---

# Catalogs

An RDBMS maintains a collection of relation instances.

To do this, it also needs information **about** relations:

- name, owner, primary key of each relation
- name, data type, constraints for each attribute
- authorisation for operations on each relation

Similarly for other DBMS objects (e.g. views, functions, triggers, ...)

This information is stored in the **system catalog**.

(The "system catalog" is also called "data dictionary" or "system view")

---

## Catalogs (cont)

DBMSs typically use a hierarchy of namespaces to manage names:

## Database (or Catalog)

- top-level namespace, contains a collection of schemas
- users typically connect to and work with a current database

## Schema

- second-level namespace, contains a collection of tables, views, etc.
- users typically work with current schema, but can qualify names

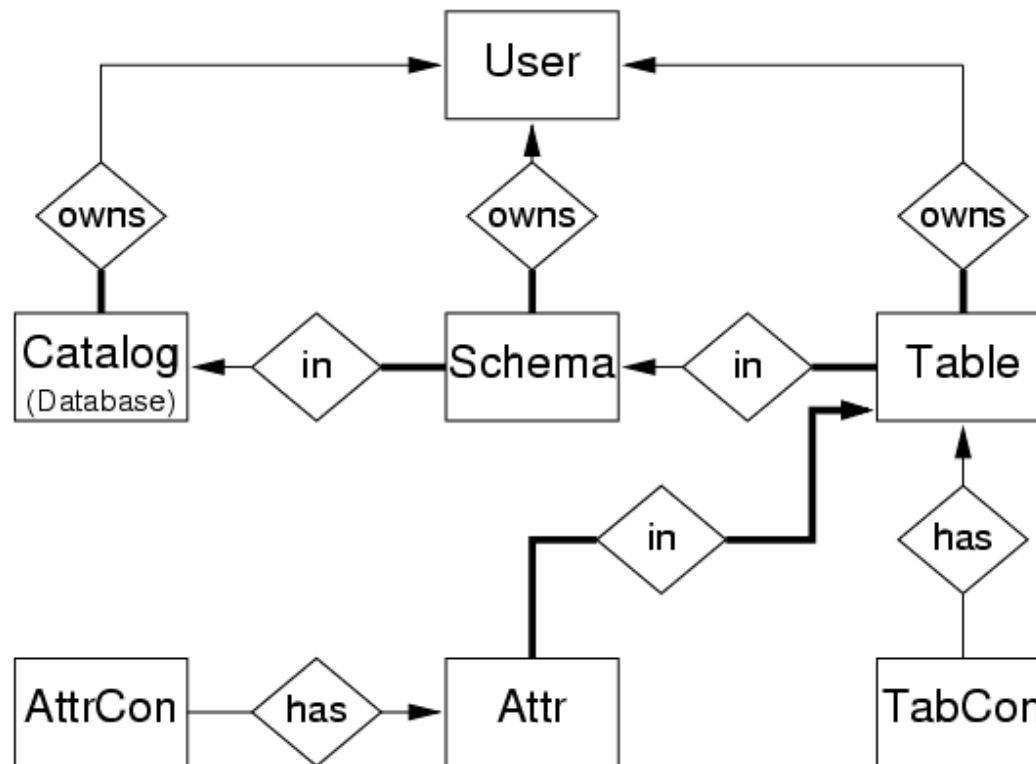
## Table

- lowest-level namespace, contains a collection of attributes
- SELECT queries set a context for names; qualification often required

---

## Catalogs (cont)

DBMSs store the catalog data in a collection of special tables:



(A small fragment of the meta-data tables in a typical RDBMS)

## Catalogs (cont)

SQL:2003 standard view of metadata: **INFORMATION\_SCHEMA**.

The **INFORMATION\_SCHEMA** is available globally and includes:

**Schemata**(catalog\_name, schema\_name, schema\_owner, ...)

**Tables**(table\_catalog, table\_schema, table\_name, table\_type, ...)

**Columns**(table\_catalog, table\_schema, table\_name, column\_name,  
ordinal\_position, column\_default, data\_type, ...)

**Views**(table\_catalog, table\_schema, table\_name, view\_definition, ...)

**Table\_Constraints**(..., constraint\_name, ..., constraint\_type, ...)

etc. etc. etc.

---

## Catalogs (cont)

DBMS internal meta-data is often different to standard, e.g.

```
Users(id:int, name:string, ...)
Databases(id:int, name:string, owner:ref(User), ...)
Schemas(id:int, name:string, owner:ref(User), ...)
Types(id:int, name:string, defn:string, size:int, ...)
Tables(id:int, name:string, owner:ref(User),
       inSchema:ref(Schema), ...)
Attributes(id:int, name:string, table:ref(Table),
           type:ref(Type), pkey:bool, ...)
```



```
TableConstraints(id:int, name:string, table:ref(Table),
                 defn:string, ...)
AttrConstraints(id:int, name:string, attr:ref(Attribute),
                defn:string, ...)
-- etc. etc. etc.
```

---

## Catalogs (cont)

SQL DDL operations such as

```
create table Abc (
    x integer primary key,
    y integer);
```

are implemented internally as operations on meta-data, e.g.

```
userID := current_user();
schemaID := current_schema();
tabID := nextval('tab_id_seq');
select into intID id
from Types where name='integer';
insert into Tables(id,name,owner,inSchema,...)
    values (tabID, 'abc', userID, schema, ...)
attrID := nextval('attr_id_seq');
```

```
insert into Attributes(id,name,table,type,pkey,...)
      values (attrID, 'x', tabID, intID, true, ...)
attrID := nextval('attr_id_seq');
insert into Attributes(id,name,table,type,pkey,...)
      values (attrID, 'y', tabID, intID, false, ...)
```

---

## Access to System Catalog

Users typically have access to the system catalog via

- special commands (e.g. PostgreSQL's `\d`, `\df`, etc.)
- query-able views (e.g. Oracle's **`select * from tab`**)

How much is visible to each user depends on their role

- DBA can see/change anything in system catalog via SQL
  - ordinary user can see only some of the system catalog and can change it only via SQL DDL statements
- 

## PostgreSQL Catalog

PostgreSQL stores catalog information as regular tables.

The `\d?` special commands in **psql** are just wrappers around queries on those tables, e.g.

**\dt** list information about tables

**\dv** list information about views

**\df** list information about functions

**\dp** list table access privileges

**\dT** list information about data types

**\dd** shows comments attached to DB objects

---

## PostgreSQL Catalog (cont)

A PostgreSQL installation typically has several databases.

Some catalog information is global, e.g.

- databases, users, ...
- there is one copy of each "global" table for the whole PostgreSQL installation
- this copy is shared by all databases in the installation

Other catalog information is local to each database, e.g

- schemas, tables, attributes, functions, types, ...
- there is a separate copy of each "local" table in each database
- a copy of each "local" table is made when a new database is created

---

## PostgreSQL Catalog (cont)

**pg\_authid** contains information about database users:

<b>oid</b>	integer key to reference user
<b>rolname</b>	symbolic user name (e.g. <b>jas</b> )
<b>rolpasswd</b>	md5–encrypted password
<b>rolcreatedb</b>	can create new databases

**rolsuper** is a superuser (owns server process)  
**rolcatupdate** can update system catalogs

---

## PostgreSQL Catalog (cont)

**pg\_database** contains information about databases:

**datname** database name (e.g. **nssis**)  
**datdba** database owner (refs **pg\_auth.oid**)  
**datpath** where files for database are stored  
(if not in the PG\_DATA directory)  
**datacl** access permissions

---

## PostgreSQL Catalog (cont)

**pg\_class** contains information about tables:

<b>relname</b>	name of table (e.g. <b>employee</b> )
<b>relnamespace</b>	schema in which table defined (refs <b>pg_namespace.oid</b> )
<b>reltype</b>	data type corresponding to table (refs <b>pg_type.oid</b> )
<b>relowner</b>	owner (refs <b>pg_authid.oid</b> )
<b>reltuples</b>	# tuples in table
<b>relacl</b>	access permissions

Also holds info about objects other than tables, e.g. views, sequences, indexes.

---

## PostgreSQL Catalog (cont)

**pg\_class** also holds various flags/counters for each table:

<b>relkind</b>	what kind of object 'r' = ordinary table, 'i' = index, 'v' = view 'c' = composite type, 'S' = sequence, 's' = special
----------------	---

<b>relnatts</b>	# attributes in table (how many entries in <b>pg_attribute</b> table)
<b>relchecks</b>	# of constraints on table (how many entries in <b>pg_constraint</b> table)
<b>relhasindex</b>	table has/had an index?
<b>relhaspkey</b>	table has/had a primary key?

etc.

---

## PostgreSQL Catalog (cont)

**pg\_type** contains information about data types:

<b>typname</b>	name of type (e.g. <b>integer</b> )
<b>typnamespace</b>	schema in which type defined (refs <b>pg_namespace.oid</b> )
<b>typowner</b>	owner (refs <b>pg_authid.oid</b> )
<b>typtype</b>	what kind of data type

'b' = base type, 'c' = complex (row) type, ...

**typelen**            how much storage used for type values  
(-1 for variable-length types, e.g. **text**)

**typrelid**           table associated to complex type  
(refs **pg\_class.oid**)

etc.

---

## PostgreSQL Catalog (cont)

**pg\_attribute** contains information about attributes:

**attname**        name of attribute (e.g. **id**)

**attrelid**       table this attribute belongs to  
(refs **pg\_class.oid**)

**atttypid**       data type of this attribute  
(refs **pg\_type.oid**)

**attlen**        storage space required by attribute



(a copy of **pg\_type.typlen** for data type)

**attnum** attribute position (1..n, sys attrs are -ve)

---

## PostgreSQL Catalog (cont)

**pg\_attribute** also holds constraint/status information:

**attnotnull** attribute may not be null?

**atthasdef** attribute has a default values  
(value is held in **pg\_attrdef** table)

**attisdropped** attribute has been dropped from table

plus others related to strange properties of attribute.

---

## PostgreSQL Catalog (cont)

**pg\_proc** contains information about functions:

<b>proname</b>	name of function (e.g. <b>substr</b> )
<b>pronamespace</b>	schema in which function defined (refs <b>pg_namespace.oid</b> )
<b>proowner</b>	owner (refs <b>pg_authid.oid</b> )
<b>prolang</b>	what language function written in
<b>proacl</b>	access control

etc.

---

## PostgreSQL Catalog (cont)

**pg\_proc** contains information about arguments:

<b>pronargs</b>	how many arguments
<b>prorettype</b>	return type (refs <b>pg_type.oid</b> )
<b>proargtypes</b>	argument types (vector of refs <b>pg_type.oid</b> )
<b>proisstrict</b>	returns null if any arg is null

**prosrc** source code if interpreted (e.g. PLpgSQL)

---

## PostgreSQL Catalog (cont)

**pg\_constraints** contains information about constraints:

<b>conname</b>	name of constraint (not unique)
<b>connamespace</b>	schema containing this constraint
<b>contype</b>	kind of constraint 'c' = check, 'u' = unique, 'p' = primary key, 'f' = foreign key
<b>conrelid</b>	which table (refs <b>pg_class.oid</b> )
<b>conkey</b>	which attributes (vector of values from <b>pg_attribute.attnum</b> )
<b>consrc</b>	check constraint expression

---

## PostgreSQL Catalog (cont)

For full details on PostgreSQL catalogs:

PostgreSQL 9.0.3 Developer's Guide

Part VII, Chapter 45, System Catalogs

Part IV, Chapter 34, The Information Schema

---

## § Security, Privilege, Authorisation

---

### Database Security

Database security has to meet the following objectives:

- **Secrecy**: information not disclosed to unauthorised users  
e.g. a student should **not** be able to examine other students' marks
- **Integrity**: only authorised users are allowed to modify data  
e.g. a student should not be able to modify anybody's marks

- **Availability**: authorised users should not be denied access  
e.g. the LIC should be able to read/changes marks for their course

Goal: prevent unauthorised use/alteration/destruction of mission–critical data.

---

## Database Security (cont)

Security mechanisms operate at a number of levels

- within the database system (SQL–level privileges)  
e.g. specific users can query/modify/update only specified database objects
- accessing the database system (users/passwords)  
e.g. users are required to authenticate themselves at connection–time
- operating system (access to DB clients)  
e.g. users should not obtain access to the DBMS superuser account
- network (most DB access nowadays is via network)  
e.g. results should not be transmitted unencrypted to Web browsers
- human/physical (conventional security mechanisms)  
e.g. no unauthorised physical access to server hosting the DBMS

## Database Access Control

Access to DBMSs involves two aspects:

- having execute permission for a DBMS client (e.g. **psql**)
- having a username/password registered in the DBMS

Establishing a **connection** to the database:

- user supplies **database/username/password** to client
- client passes these to server, which validates them
- if valid, user is "logged in" to the specified database

---

## Database Access Control (cont)

Note: we don't need to supply username/password to **psql**

- **psql** works out which user by who ran the client process
- we're all PostgreSQL super-users on our own servers

- servers are configured to allow super-user direct access

Note: access to databases via the Web involves:

- running a script on a Web server
- using the Web server's access rights on the DBMS

---

## Database Access Control (cont)

Database **users** are set up by the DBA via the command:

```
CREATE USER Name IDENTIFIED BY 'Password'
```

Various privileges can be assigned at user-creation time, e.g.

- ability to create new users or databases or tables

User properties may be subsequently changed via:

```
ALTER USER Name IDENTIFIED BY 'NewPassword'
```

This command is also used to change privileges, quotas, etc.

---

## Database Access Control (cont)

A user may be associated with a **role** or **group**

- which typically gives them additional specific privileges

Roles are also set up by the DBA via the command:

```
CREATE ROLE RoleName
```

Examples of roles:

- AcademicStaff ... has privileges to read/modify marks
  - OfficeStaff ... has privilege to read all marks
  - Student ... has privilege to read own marks only
- 

## Database Access Control in PostgreSQL



In older versions of PostgreSQL ...

- **USERS** and **GROUPS** were distinct kinds of objects
- **USERS** were added via **CREATE USER** *UserName*
- **GROUPS** were added via **CREATE GROUP** *GroupName*
- **GROUPS** were built via **ALTER GROUP ... ADD USER ...**

In recent versions, **USERS** and **GROUPS** are unified by **ROLES**

Older syntax is retained for backward compatibility.

---

## Database Access Control in PostgreSQL (cont)

PostgreSQL has two ways to create users ...

From the Unix command line, via the command

```
createuser [ -a | -d ] Name
```

(**-a** allows user to create other users, **-d** allows user to create databases)

From SQL, via the statement:

```
CREATE ROLE UserName Options
-- where Options include ...
PASSWORD 'Password'
CREATEDB | NOCREATEDB
CREATEUSER | NOCREATEUSER
IN GROUP GroupName
VALID UNTIL 'TimeStamp'
```

---

## Database Access Control in PostgreSQL (cont)

Groups are created as **ROLES** via

```
CREATE ROLE GroupName
--or--
CREATE ROLE GroupName
WITH USER UserName1, UserName2, ...
```

and may be subsequently modified by

```
GRANT GroupName TO UserName1, UserName2, ...  
--and  
REVOKE GroupName FROM UserName1, UserName2, ...
```

---

## Database Access Control in PostgreSQL (cont)

PostgreSQL stores user information in a table **pg\_authid**.

Each user is associated with a unique identifying number.

Every PostgreSQL is created with a default user (id=1, superuser).

Some fields in the **pg\_authid** table:

- **oid**: unique user id (e.g. 1, 100, ...)
  - **rolname**: string for user name (e.g. **jas**)
  - **rolpassword**: encrypted version of user's password
- 

## Database Access Control in PostgreSQL (cont)

PostgreSQL uses a file called **pg\_hba.conf** to determine how to authenticate users (**hba** stands for host-based authentication).

This file contains a sequence of entries where each entry has:

- connection method (Unix-domain socket, TCP/IP, encrypted)
- list of databases (or **all**)
- user/group name (or **all**)
- IP address and IP mask
- authentication method (plus options)

PostgreSQL uses first four items to determine authentication method.

---

## Database Access Control in PostgreSQL (cont)

Possible authentication methods:

- **md5**: get password from user, check against **pg\_shadow**
- **krb5**: use Kerberos-based authentication

- **reject**: do not allow the user to access DBMS this way
  - **trust**: allow them to log in as any user without password
  - **ident**: allow them access based on Unix user name info
    - **ident+map**: see if UnixName is member of map list
    - **ident+sameuser**: log in as user with name UnixName
- 

## Database Access Control in PostgreSQL (cont)

Examples from **pg\_hba.conf** file:

```
# Allow any user on the local system to connect to any database
# as any user if accessing from local machine via Unix socket.
#
```

```
#TYPE    DATABASE USER IP-ADDRESS      IP-MASK          METHOD
local   all      all                               trust
```

```
# Reject all connection from 192.168.54.1
#
```

```
#TYPE    DATABASE USER IP-ADDRESS      IP-MASK          METHOD
host    all      all 192.168.54.1 255.255.255.255 reject
```

```
# Allow any user from any host with IP address 192.168.93.x
```

```
# to connect to database "mydb" as the same user name that
# ident reports for the connection (typically Unix user name).
#
#TYPE  DATABASE USER IP-ADDRESS      IP-MASK          METHOD
host  mydb    all  192.168.93.0  255.255.255.0  ident sameuser

# Allow a user from host 192.168.12.10 to connect to database
# "mydb" only if the user's password is correctly supplied.
#
#TYPE  DATABASE USER IP-ADDRESS      IP-MASK          METHOD
host  mydb    all  192.168.12.10 255.255.255.255 md5
```

---

## SQL Access Control

SQL access control deals with

- privileges on database objects (e.g. tables, view, functions, ...)
- allocating such privileges to users and/or roles

The user who creates an object is automatically assigned:

- ownership of that object
- a privilege to modify (**ALTER**) the object

- a privilege to remove (**DROP**) the object
  - along with all other privileges specified below
- 

## SQL Access Control (cont)

The owner of an object can assign privileges on that object to other users.

Accomplished via the command:

```
GRANT Privileges ON Object  
TO list of (Users|Roles) | PUBLIC  
[ WITH GRANT OPTION ]
```

*Privileges* can be **ALL** (giving everything but **ALTER** and **DROP**)

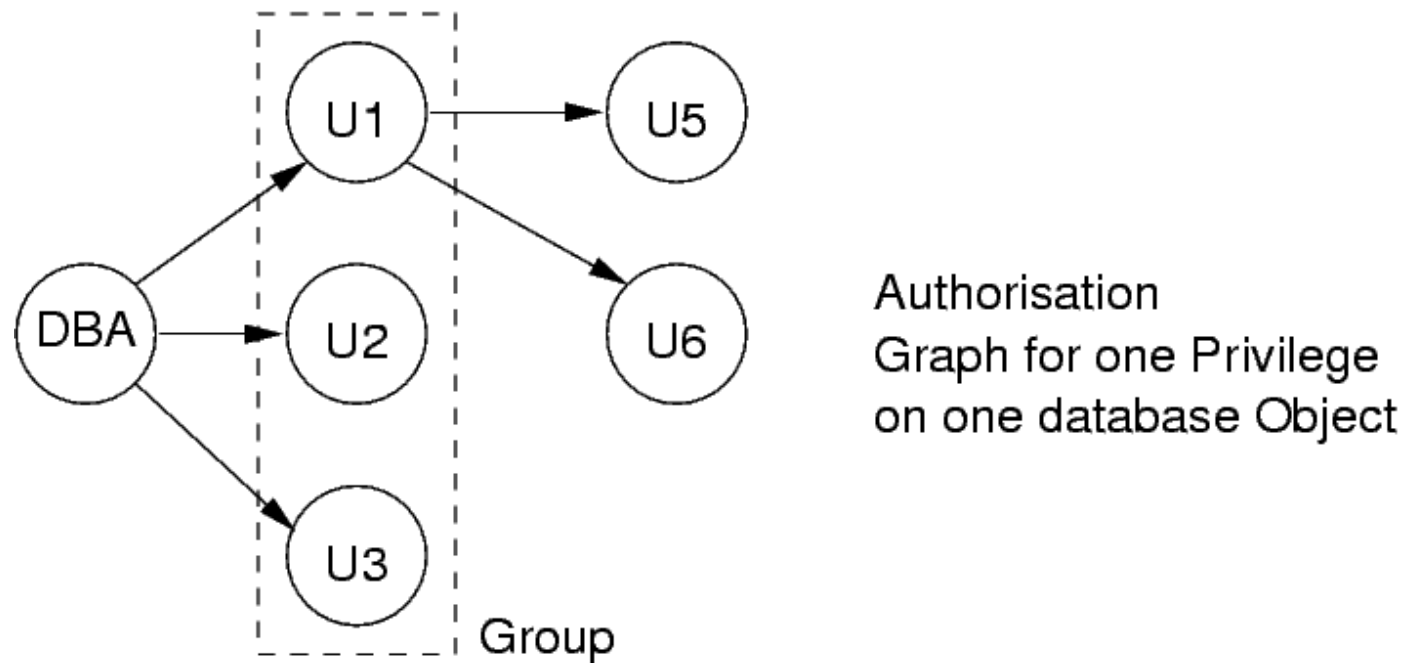
**WITH GRANT OPTION** allows a user who has been granted a privilege to pass the privilege on to any other user.

---

## SQL Access Control (cont)

## Effects of privilege granting

- are sometimes subtle (possible conflicts?)
- can be represented by an **authorisation graph**



---

## SQL Access Control (cont)

Privileges can be withdrawn via the command:



```
REVOKE Privileges ON Object  
FROM ListOf (Users|Roles) | PUBLIC  
CASCADE | RESTRICT
```

Normally withdraws Privileges from just specified users/roles.

**CASCADE** ... also withdraws from users they had granted to.

E.g. revoking from U1 also revokes U5 and U6

**RESTRICT** ... fails if users had granted privileges to others.

E.g. revoking from U1 fails, revoking U5 or U2 succeeds

---

## SQL Access Control (cont)

Privileges available for users on database objects:

**SELECT:**

- user can read all rows and columns of table/view

- this includes columns added later via **ALTER TABLE**

**INSERT** or **INSERT**(*ColName*):

- user can insert rows into table
  - if *ColName* specified, can only set value of that column
- 

## SQL Access Control (cont)

More privileges available for users on database objects:

**UPDATE** or **UPDATE**(*ColName*):

- user can modify values stored in the table
- if *ColName* specified, can only set value of that column

**DELETE:**

- user can delete rows from the table
- this does *not* imply permission to remove table itself  
(this is the **DROP** privilege automatically assigned to the object creator)

## SQL Access Control (cont)

More privileges available for users on database objects:

### **REFERENCES** (*ColName*):

- user can use *ColName* as foreign key in their tables

### **EXECUTE**:

- user can execute the specified function

### **TRIGGER**:

- user is allowed to create triggers on a table
  - note that triggers always execute with creator's privileges
- 

## SQL Access Control in PostgreSQL

PostgreSQL follows the above with some minor variations:

- group names must be preceded by **GROUP**
- there is an additional privilege for **RULES**

See the PostgreSQL manual for full details.

---

## Problems with SQL Access Control

Allowing users to assign privileges to others can be exploited.

Example: student S wants access to table of marks M

- M is owned by lecturer L, S has no SELECT privilege on M
  - S creates a new table SS and grants INSERT privilege to L
  - S modifies code of some PLpgSQL function F used by L
  - the modifications to F copy the data from M into SS
  - S restores F to its original state (in case L gets suspicious)
  - S now has access to the data from M in table SS
-

# Mandatory Access Control

Above approach is called **discretionary access control**

- relies on individual users to assign privileges
- very fine-grained  $\Rightarrow$  tedious to specify privileges

An alternative approach: **mandatory access control** (MAC)

- global approach to access control; simple to specify
- currently under development (not yet available in DBMSs)
- a popular MAC model is Bell–LaPadula (see next page)

---

## Mandatory Access Control (cont)

Access control is described in terms of ...

- **objects** (e.g. tables, rows, views, columns, ...)
- **subjects** (e.g. users, programs, ...)

- **security classes** (an ordered collection, least to most secure)

Example: *Unclassified* < *Confidential* < *Secret* < *Top Secret*

Subjects and objects are assigned to security classes.

Impose these restrictions on every data access:

- subject S can read object O only if  $\text{class}(S) \geq \text{class}(O)$
- subject S can write object O only if  $\text{class}(S) \leq \text{class}(O)$

---

## Security in Statistical Databases

### Statistical databases

- contain data about a group of individuals
- but allow access only to summary data
- also provide controls to select subsets of data

Example of such a database: population census

- information is stored about individuals
- users are only allowed to examine trends/summaries
- e.g. can find out average age of people living in Sydney but cannot ask the question "How old is John?"

Privacy is protected, but useful information is still available.

---

## Security in Statistical Databases (cont)

Consider: anonymous surveys in an on-line survey system

- can see overall results (e.g. tutorials rated at 7/10)
- but do not have access to individual student responses
- can also summarize results for subgroups  
(e.g. CE students rated tutorials at 6/10, CS students rated tutorials at 8/10)

Problem: subset controls may allow selection of individuals, e.g.

- we know that there's only one Law student in the class

- ask for a "summary" of responses given by Law students
- 

## Security in Statistical Databases (cont)

How to solve this problem?

- restrict subsets to being larger than a minimum size  $N$

But still doesn't quite work, e.g.

- system gives summaries for several groups (e.g. SE+CE)
- get summary for Law+CE (presumably with  $>N$  responses)
- get summary just for CE (assume count of responses given)
- determine Law response from the "difference"

Security is difficult to enforce in statistical databases, but activity can be logged.

---

## § Performance Tuning

---



# Performance Tuning

Schema design:

- devise data structures to represent application information

Performance tuning:

- devise data structures to achieve good performance

Good performance may involve any/all of:

- making applications run faster
- lowering response time of queries/transactions
- improving overall transaction throughput

---

## Performance Tuning (cont)

Tuning requires us to consider the following:

- which queries and transactions will be used?  
(e.g. check balance for payment, display recent transaction history)
  - how frequently does each query/transaction occur?  
(e.g. 99% of transactions are EFTPOS payments; 1% are print balance)
  - are there time constraints on queries/transactions?  
(e.g. payment at EFTPOS terminals must be approved within 7 seconds)
  - are there uniqueness constraints on any attributes?  
(therefore, define index on attributes to speed up insertion uniqueness check)
  - how frequently do updates occur?  
(indexes slow down updates, because must update table *and* index)
- 

## Performance Tuning (cont)

Performance can be considered at two times:

- *during* schema design
  - typically towards the end of schema design process

- requires schema transformations such as denormalisation
  - *after* schema design
    - requires adding extra data structures such as indexes
- 

## Denormalisation

Normalisation structures data to minimise storage redundancy.

- achieves this by "breaking up" the data into logical chunks
- requires minimal "maintenance" to ensure data consistency

Problem: queries that need to put data back together.

- need to use a (potentially expensive) join operation
- if an expensive join is frequent, system performance suffers

Solution: store some data redundantly

- benefit: queries needing expensive join are now cheap
- trade-off: extra maintenance effort to maintain consistency

- worthwhile if joins are frequent and updates are rare
- 

## Denormalisation (cont)

Example: Courses = Course  Subject  Term

If we frequently need to refer to course "standard" name

- add extra **courseName** column into **Course** table
- cost: trigger before insert on **Course** to construct name
- trade-off likely to be worthwhile: **Course** insertions infrequent

```
-- can now replace a query like:
select s.code||t.year||t.sess, e.grade, e.mark
from   Course c, CourseEnrolment e, Subject s, Term t
where  e.course = c.id and c.subject = s.id and c.term = t.id
-- by a query like:
select c.courseName, e.grade, e.mark
from   Course c, CourseEnrolment e
where  e.course = c.id
```

---

# Indexes

Indexes provide efficient content-based access to tuples.

Can build indexes on any (combination of) attributes.

Defining indexes:

```
CREATE INDEX name ON table ( attr1, attr2, ... )
```

*attr*<sub>*i*</sub> can be an arbitrary expression (e.g. **upper(name)**).

**CREATE INDEX** also allows us to specify

- that the index is on **UNIQUE** values
- an access method (**USING** btree, hash, rtree, or gist)

---

## Indexes (cont)

Indexes can make a huge difference to query processing cost.

On the other hand, they introduce overheads (storage, updates).

Creating indexes to maximise performance benefits:

- apply to attributes used in equality/range conditions, e.g.

```
select * from Employee where id = 12345
select * from Employee where age > 60
select * from Employee where salary between 10000 and 20000
```

- but only in queries that are frequently used
  - and on tables that are not updated frequently
- 

## Indexes (cont)

Considerations in applying indexes:

- is an attribute used in frequent/expensive queries?  
(note that some kinds of queries can be answered from index alone)

- should we create an index on a collection of attributes?  
(yes, if the collection is used in a frequent/expensive query)
- can we exploit a clustered index? (only one per table)
- should we use B-tree or Hash index?

```
-- use hashing for (unique) attributes in equality tests, e.g.  
select * from Employee where id = 12345  
-- use B-tree for attributes in range tests, e.g.  
select * from Employee where age > 60
```

---

## Query Tuning

Sometimes, a query can be re-phrased to affect performance:

- by helping the optimiser to make use of indexes
- by avoiding (unnecessary) operations that are expensive

Examples which *may* prevent optimiser from using indexes:

```
select name from Employee where salary/365 > 10.0
    -- fix by re-phrasing condition to (salary > 3650)
select name from Employee where name like '%ith%'
select name from Employee where birthday is null
    -- above two are difficult to "fix"
select name from Employee
where dept in (select id from Dept where ...)
    -- fix by using Employee join Dept on (e.dept=d.id)
```

---

## Query Tuning (cont)

Other factors to consider in query tuning:

- **select distinct** requires a sort; is **distinct** necessary?
- if multiple join conditions are available ...  
choose join attributes that are indexed, avoid joins on strings

```
select ... Employee join Customer on (s.name = p.name)
vs
select ... Employee join Customer on (s.ssn = p.ssn)
```



- sometimes **or** in condition prevents index from being used ...  
replace the **or** condition by a union of non-**or** clauses

```
select name from Employee where dept=1 or dept=2  
vs  
(select name from Employee where dept=1)  
union  
(select name from Employee where dept=2)
```

---

## PostgreSQL Query Tuning

PostgreSQL provides the **explain** statement to

- give a representation of the query execution plan
- with information that may help to tune query performance

Usage:

```
EXPLAIN [ANALYZE] Query
```

Without **ANALYZE**, **EXPLAIN** shows plan with estimated costs.

With **ANALYZE**, **EXPLAIN** executes query and prints real costs.

Note that runtimes may show considerable variation due to buffering.

---

## EXPLAIN Examples

Example: Select on indexed attribute

```
ass2=# explain select * from student where id=100250;  
                QUERY PLAN
```

```
-----  
Index Scan using student_pkey on student  (cost=0.00..5.94 rows=1 width=17)  
Index Cond: (id = 100250)
```

```
ass2=# explain analyze select * from student where id=100250;  
                QUERY PLAN
```

```
-----  
Index Scan using student_pkey on student  (cost=0.00..5.94 rows=1 width=17)  
                                           (actual time=31.209..31.212 rows=1 loops=1)  
Index Cond: (id = 100250)  
Total runtime: 31.252 ms
```

---

## EXPLAIN Examples (cont)

## Example: Select on non-indexed attribute

```
ass2=# explain select * from student where stype='local';  
               QUERY PLAN
```

```
-----  
Seq Scan on student  (cost=0.00..70.33 rows=18 width=17)  
  Filter: ((stype)::text = 'local'::text)
```

```
ass2=# explain analyze select * from student where stype='local';  
               QUERY PLAN
```

```
-----  
Seq Scan on student  (cost=0.00..70.33 rows=18 width=17)  
    (actual time=0.061..4.784 rows=2512 loops=1)  
  Filter: ((stype)::text = 'local'::text)  
Total runtime: 7.554 ms
```

---

## EXPLAIN Examples (cont)

### Example: Join on a primary key (indexed) attribute

```
ass2=# explain  
ass2=# select s.sid,p.name from Student s, Person p where s.id=p.id;  
               QUERY PLAN
```

---

```

Hash Join  (cost=70.33..305.86 rows=3626 width=52)
  Hash Cond: ("outer".id = "inner".id)
    -> Seq Scan on person p  (cost=0.00..153.01 rows=3701 width=52)
    -> Hash  (cost=61.26..61.26 rows=3626 width=8)
        -> Seq Scan on student s  (cost=0.00..61.26 rows=3626 width=8)

```

---

## EXPLAIN Examples (cont)

Join on a primary key (indexed) attribute:

```

ass2=# explain analyze
ass2=# select s.sid,p.name from Student s, Person p where s.id=p.id;
               QUERY PLAN

```

---

```

Hash Join  (cost=70.33..305.86 rows=3626 width=52)
    (actual time=11.680..28.242 rows=3626 loops=1)
  Hash Cond: ("outer".id = "inner".id)
    -> Seq Scan on person p  (cost=0.00..153.01 rows=3701 width=52)
        (actual time=0.039..5.976 rows=3701 loops=1)
    -> Hash  (cost=61.26..61.26 rows=3626 width=8)
        (actual time=11.615..11.615 rows=3626 loops=1)
        -> Seq Scan on student s  (cost=0.00..61.26 rows=3626 width=8)
            (actual time=0.005..5.731 rows=3626 loops=1)

Total runtime: 32.374 ms

```

---

Produced: 13 Sep 2020