

Mapping ER to SQL

- Mapping ER to SQL
- Reminder: SQL/Relational Model vs ER Model
- Mapping ER to SQL
- Mapping Strong Entities
- Mapping Weak Entities
- Mapping N:M Relationships
- Mapping 1:N Relationships
- Mapping 1:1 Relationships
- Mapping n-way Relationships
- Mapping Composite Attributes
- Mapping Multi-valued Attributes (MVAs)
- Mapping Subclasses



❖ Mapping ER to SQL

We have explored mapping ER designs to relational schemas

SQL schemas are essentially more detailed versions of relational schemas

The mapping is much the same, except that

- you need to provide more details on allowed values
- you can map some ideas from ER that are not in relational schemas

There are also some ideas from ER that do not map to an SQL schema

❖ Reminder: SQL/Relational Model vs ER Model

Correspondences between SQL/relational and ER data models:

- $\text{attribute(ER)} \cong \text{attribute(Rel)}$, $\text{entity(ER)} \cong \text{row/tuple(Rel)}$
- $\text{entity set(ER)} \cong \text{table/relation(Rel)}$, $\text{relationship(ER)} \cong \text{table/relation(Rel)}$

Differences between SQL and ER models:

- SQL uses tables to model entities *and* relationships
- SQL has no composite or multi-valued attributes (only atomic)
- SQL has no object-oriented notions (e.g. subclasses, inheritance)

Note that ...

- not all aspects of ER can be represented exactly in an SQL schema
- some aspects of SQL schemas (e.g. domains) do not appear in ER

❖ Mapping ER to SQL

Some conventions that we use in mapping ER to SQL

- stop using upper-case for SQL keywords (use **table** vs **TABLE**)
- all tables based on entities are given plural names
- attributes in entities are given the same name in ER and SQL
- attributes in relationships are given the same name in ER and SQL
- ER key attributes are defined using **primary key**
- text-based attributes are defined with type **text**, unless there is a size which is obvious from the context
- attribute domains can be PostgreSQL-specific types where useful
- foreign keys within entity tables are named after the relationship
- foreign keys in relationship tables are named **table_id**


❖ Mapping Strong Entities

An entity set E with atomic attributes a_1, a_2, \dots, a_n

maps to

A table R with attributes (columns) a_1, a_2, \dots, a_n

Example:



SQL Version


```
create table Students (  
    stuID integer primary key,  
    name text not null,  
    degree char(4)  
);
```

Note: the key is preserved in the mapping.

❖ Mapping Weak Entities

Example:

ER Model



SQL Version

```


create table Employees (
  SSN    text primary key,
  ename  text,
  salary currency
);
  
```

```

create table Contacts (
  relatedTo text not null, -- total participation
  name      text,         -- not null implied by PK
  phone     text not null,
  primary key (relatedTo, name),
  foreign key (relatedTo) references Employees (ssn)
);
  
```

❖ Mapping N:M Relationships

Example:



Relational Version

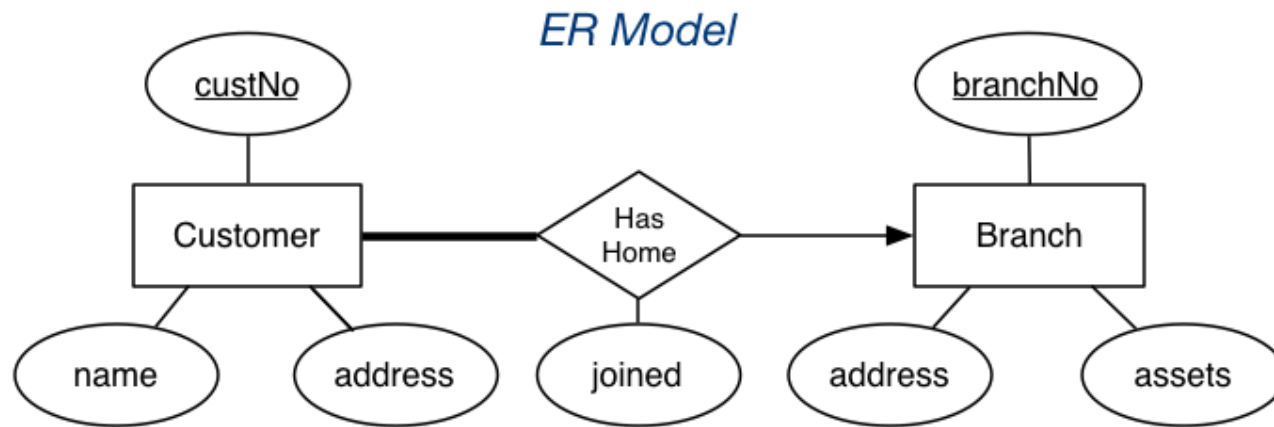
Customer	custNo	name	address
Account	acctNo	title	balance
Owns	acctNo	custNo	lastAccessed

❖ Mapping N:M Relationships (cont)

```
create table Customers (  
    custNo  serial primary key,  
    name    text not null,  
    address text -- don't need to know customer's address  
);  
create table Accounts (  
    acctNo  char(5) check (acctNo ~ '[A-Z]-[0-9]{3}'),  
    title   text not null,      -- acctNos are like 'A-123'  
    balance float default 0.0,  
    primary key (acctNo)  
);  
create table Owns (  
    customer_id integer references Customers(custNo),  
    account_id   char(5) references Accounts(acctNo),  
    last_accessed timestamp,  
    primary key (customer_id, account_id)  
);
```


❖ Mapping 1:N Relationships

Example:



Relational Version

Customer	custNo	name	address	branchNo	joined
----------	---------------	------	---------	----------	--------

Branch	branchNo	address	assets
--------	-----------------	---------	--------


❖ Mapping 1:N Relationships (cont)

```
create table Branches (  
    branchNo serial primary key,  
    address  text not null,  
    assets   currency  
);  
create table Customers (  
    custNo  serial primary key,  
    name    text not null,  
    address text,  
    hasHome integer not null, -- total participation  
    joined  date not null,  
    foreign key (hasHome) references Branches(branchNo)  
);
```

hasHome implements the 1:n relationship; **not null** implements total participation

❖ Mapping 1:1 Relationships

Example:



Relational Version

Manager	empNo	name	salary	branchNo
---------	--------------	------	--------	----------

Branch	branchNo	address	assets
--------	-----------------	---------	--------

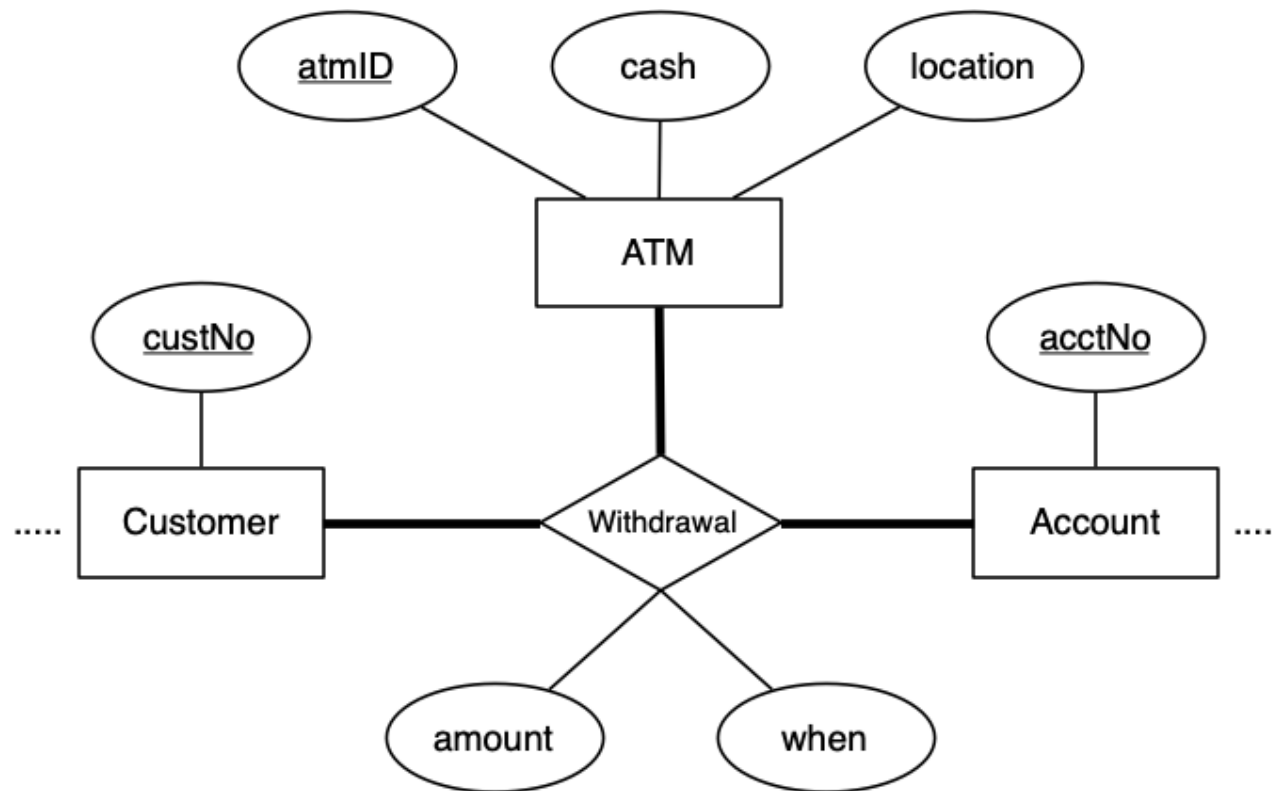
❖ Mapping 1:1 Relationships (cont)

```
create table Branches (  
    branchNo serial primary key,  
    address  text not null,  
    assets   currency          -- a new branch  
);                               -- may have no accounts  
create table Managers (  
    empNo    serial primary key,  
    name     text not null,  
    salary   currency not null, -- when first employed,  
                                -- must have a salary  
    manages  integer not null, -- total participation  
    foreign key (manages) references Branches(branchNo)  
);
```

If both entities have total participation, cannot express this in SQL except by putting a (redundant) **not null** foreign key in one table

❖ Mapping n-way Relationships

Example:



A customer accesses one of their accounts at a specific ATM


❖ Mapping n-way Relationships (cont)

```
create table Customers (  
    custNo    serial primary key, ...  
);  
create table Accounts (  
    acctNo    char(5) ... primary key, ...  
);  
create table ATMs (  
    atmID     serial primary key,  
    cash      currency check (cash >= 0),  
    location  text not null  
);  
create table Withdrawal (  
    customer_id integer references Customers(custNo),  
    account_id   char(5) references Accounts(acctNo),  
    atm_id       integer references ATMs(atmID),  
    amount       currency not null,  
    when         timestamp default now(),  
    primary key  (customer_id, account_id, atm_id)  
);
```

❖ Mapping Composite Attributes

Composite attributes are mapped by concatenation or flattening.

Example:



Relational Version #1

Person

SSN	name	address
------------	------	---------

Relational Version #2

Person

SSN	given	family	
.....	number	street	suburb	pcode

❖ Mapping Composite Attributes (cont)

-- Version 1: concatenated

```
create table People (  
    ssn      integer primary key,  
    name     text not null,  
    address  text not null  
);
```

-- Version 2: flattened

```
create table People (  
    ssn      integer primary key,  
    given    text not null,  
    family   text,  
    number   integer not null,  
    street   text not null,  
    suburb   text not null,  
    pcode    char(4) not null check (pcode ~ '[0-9]{4}')
```

```
);  
  
address = (number::text || ' ' || street || ', ' || suburb || ' ' || pcode)
```

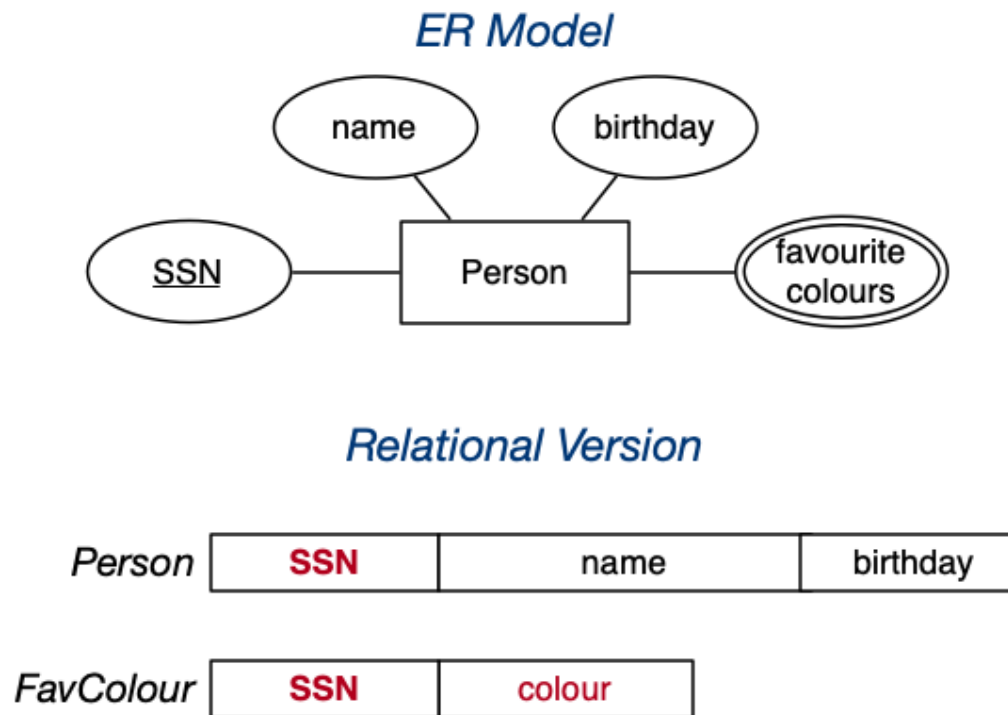
Searching: **suburb** = 'Coogee' vs **address like** '%Coogee%'

Sorting: **order by family** vs can't be done (easily)

❖ Mapping Multi-valued Attributes (MVAs)

MVAs are mapped by a new table linking values to their entity.

Example:



❖ Mapping Multi-valued Attributes (MVAs) (cont)

```
create table People (  
    ssn          integer primary key,  
    name         text not null,  
    birthday     date  
);  
create table FavColour (  
    person_id integer references People(ssn),  
    colour      text,  
    primary key (person_id, colour)  
);
```

Note that **colour** is implicitly **not null** because it is part of the primary key

❖ Mapping Subclasses


Three different approaches to mapping subclasses to tables:

- ER style
 - each entity becomes a separate table,
 - containing attributes of subclass + FK to superclass table
- object-oriented
 - each entity becomes a separate table,
 - inheriting all attributes from all superclasses
- single table with nulls
 - whole class hierarchy becomes one table,
 - containing all attributes of all subclasses (null, if unused)

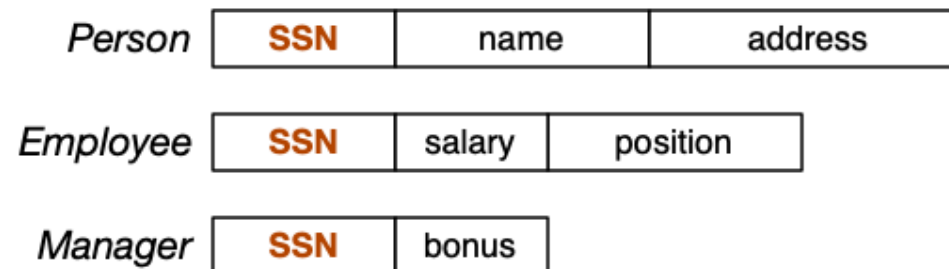
Which mapping is best depends on how data is to be used.

❖ Mapping Subclasses (cont)

Example of ER-style mapping:



Relational Version

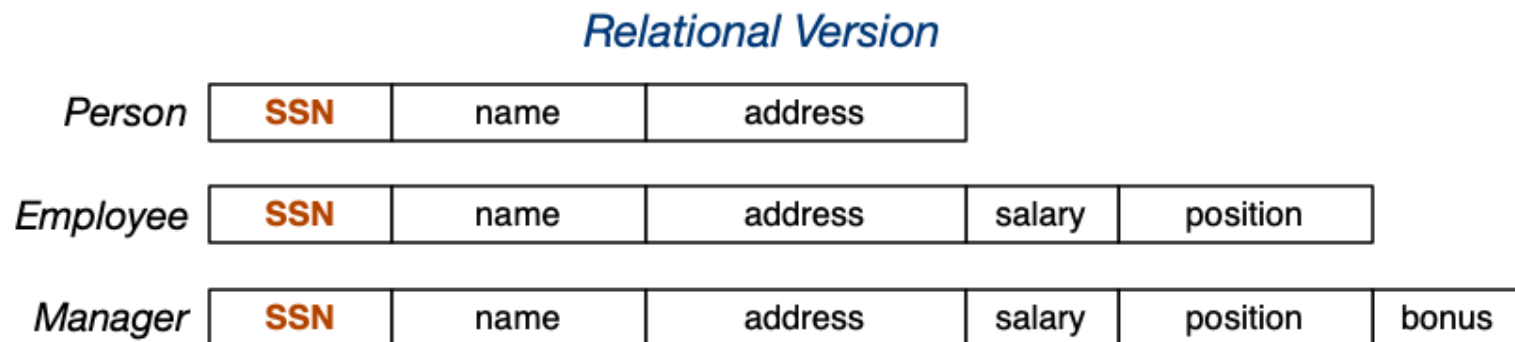
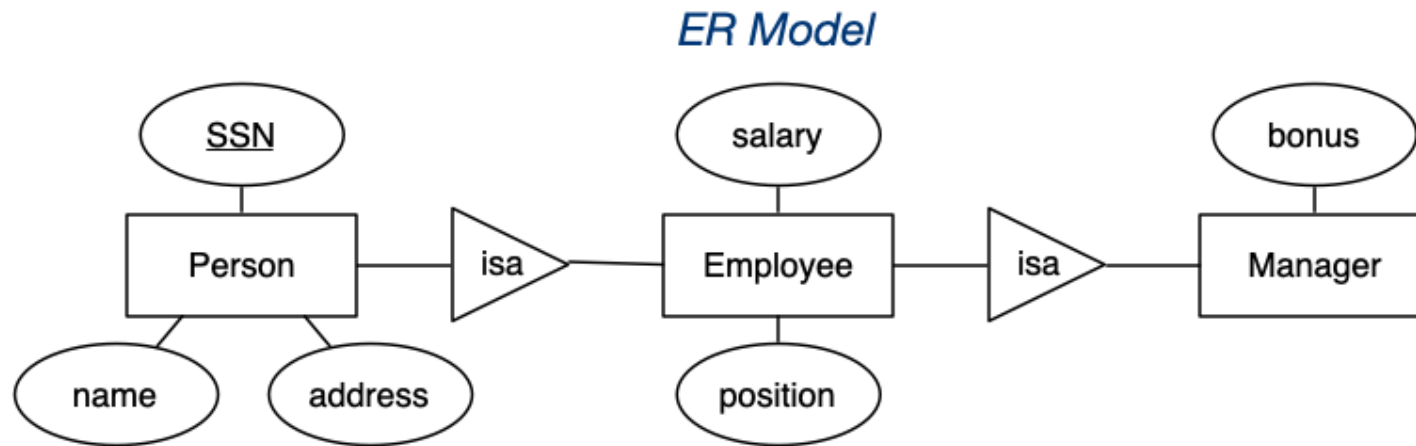


❖ Mapping Subclasses (cont)

```
create table People (  
    ssn      integer primary key,  
    name     text not null,  
    address  text  
);  
create table Employees (  
    person_id integer primary key,  
    salary     currency not null,  
    position  text not null,  
    foreign key (person_id) references People(ssn)  
);  
create table Managers (  
    employee_id integer primary key,  
    bonus        currency,  
    foreign key (employee_id)  
                references Employees(person_id)  
);
```

❖ Mapping Subclasses (cont)

Example of object-oriented mapping:

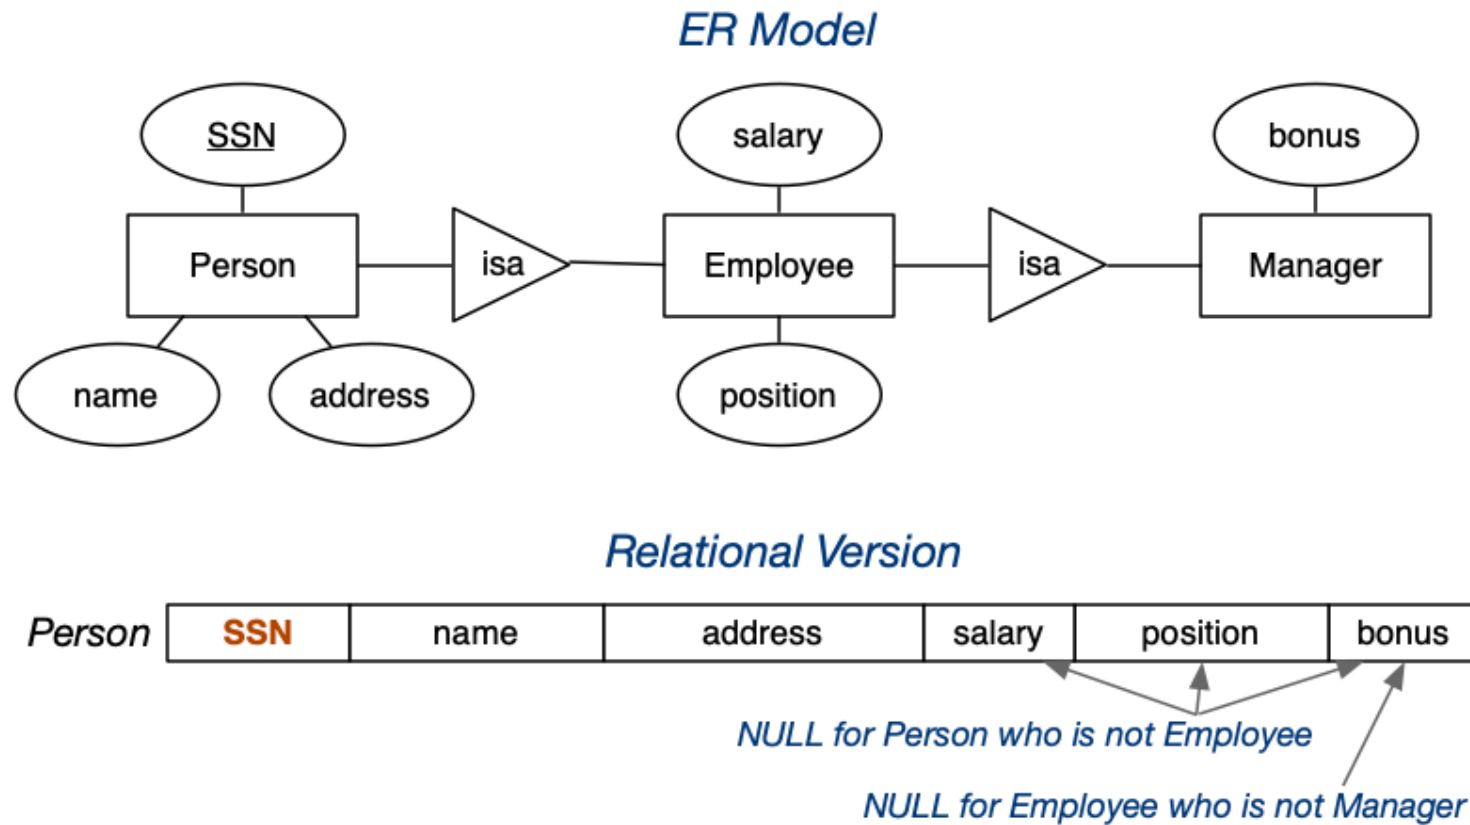


❖ Mapping Subclasses (cont)

```
create table People (  
    ssn      integer primary key,  
    name     text not null,  
    address  text  
);  
create table Employees (  
    ssn      integer primary key,  
    name     text not null,  
    address  text  
    salary   currency not null,  
    position text not null,  
    foreign key (ssn) references People(ssn)  
);  
create table Managers (  
    ssn      integer primary key,  
    name     text not null,  
    address  text  
    salary   currency not null,  
    position text not null,  
    bonus    currency,  
    foreign key (ssn) references People(ssn)  
);
```


❖ Mapping Subclasses (cont)

Example of single-table-with-nulls mapping:

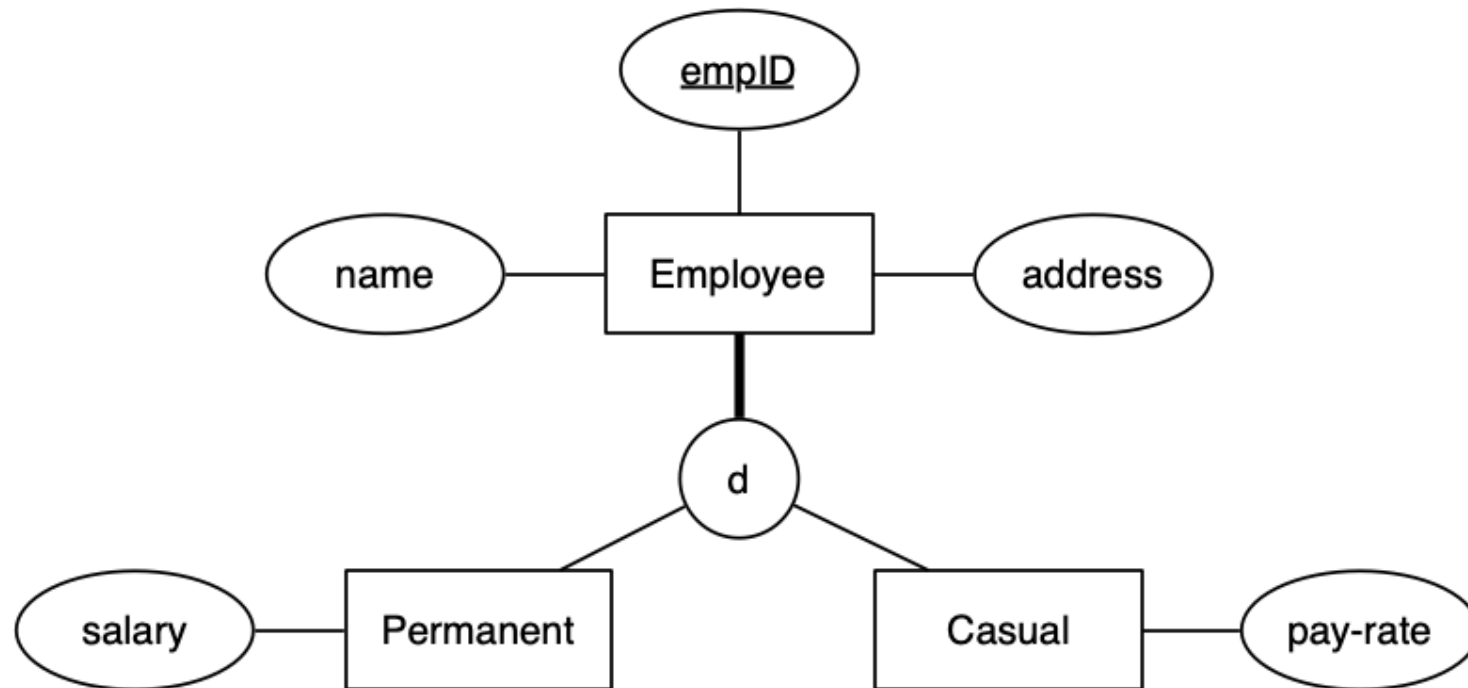


❖ Mapping Subclasses (cont)

```
create table People (  
    ssn          integer primary key,  
    ptype        char(1)  not null  
                    check (ptype in ('P','E','M')),  
    name         text not null,  
    address      text  
    salary       currency,  
    position     text,  
    bonus        currency,  
    constraint subclasses check  
        ((ptype = 'P' and salary is null  
            and position is null and bonus is null)  
        or  
        (ptype = 'E' and salary is not null  
            and position is not null and bonus is null)  
        or  
        (ptype = 'M' and salary is not null  
            and position is not null and bonus is not null))  
);
```

❖ Mapping Subclasses (cont)

Example:



Every employee is either permanent or casual, but not both.

❖ Mapping Subclasses (cont)

ER-style mapping to SQL schema:

```
create table Employees (  
    empID    serial primary key,  
    name     text not null,  
    address  text not null  
);  
create table Permanents (  
    employee_id integer primary key,  
    salary       currency not null,  
    foreign key (employee_id) references Employees(empID)  
);  
create table Casuals (  
    employee_id integer primary key,  
    pay_rate    currency not null,  
    foreign key (employee_id) references Employees(empID)  
);
```

Does *not* capture either participation or disjoint-ness constraints!

Would need to program a solution to this e.g web-form that requires user to enter both Employee and subclass info

Produced: 10 Feb 2021