

SQL Queries (iv): Grouping

- Grouping
- Restrictions on **SELECT** Lists
- Filtering Groups
- Partitions

❖ Grouping

SELECT-FROM-WHERE can be followed by **GROUP BY** to:

- partition result relation into groups (according to values of specified attribute)
- summarise (aggregate) some aspects of each group
- output one tuple per group, with grouping attribute and aggregates

R

A	B
1	'a'
2	'b'
3	'a'
1	'b'
2	'a'
1	'c'

R group by A

A	B
1	'a'
1	'b'
1	'c'
2	'b'
2	'a'
3	'a'

A, count(*), max(B)

A	count	max
1	3	'c'
2	2	'b'
3	1	'a'

❖ Grouping (cont)

Example: How many different beers does each brewer make?

```
SELECT    brewer, COUNT(name) as nbeers
FROM      Beers
GROUP BY  brewer;
```

brewer	nbeers
West City	1
James Squire	5
Yullis	1
Hop Nation	4
Anderson Valley	1
Beatnik	1
Boatrock	3
Kizakura	1
...	

❖ Grouping (cont)

GROUP BY is used as follows:

```
SELECT    attributes/aggregations
FROM      relations
WHERE     condition
GROUP BY  attributes
```

Semantics:

1. apply product and selection as for **SELECT-FROM-WHERE**
2. partition result into groups based on values of *attributes*
3. apply any aggregation separately to each group

Grouping is typically used in queries involving the phrase "for each".

❖ Restrictions on **SELECT** Lists

When using grouping, every attribute in the **SELECT** list must:

- have an aggregation operator applied to it OR
- appear in the **GROUP-BY** clause

Incorrect Example: Find the styles associated with each brewer

```
SELECT    brewer, style
FROM      Beers
GROUP BY  brewer;
```

PostgreSQL's response to this query:

```
ERROR: column beers.style must appear in the GROUP BY
       clause or be used in an aggregate function
```

❖ Filtering Groups

In some queries, you can use the **WHERE** condition to eliminate groups.

Example: Average beer price by suburb excluding hotels in The Rocks.

```
SELECT    b.addr, AVG(s.price)
FROM      Sells s join Bars b on (s.bar=b.name)
WHERE     b.addr <> 'The Rocks'
GROUP BY  b.addr;
```

For conditions on whole groups, use the **HAVING** clause.

❖ Filtering Groups (cont)

HAVING is used to qualify a **GROUP-BY** clause:

```
SELECT    attributes/aggregations
FROM      relations
WHERE     condition1      (on tuples)
GROUP BY  attributes
HAVING    condition2;    (on group)
```

Semantics of **HAVING**:

1. generate the groups as for **GROUP-BY**
2. discard groups **not** satisfying **HAVING** condition
3. apply aggregations to remaining groups

❖ Filtering Groups (cont)

Example: Number of styles from brewers who make at least 5 beers?

```
SELECT    brewer, count(name) as nbeers,  
          count(distinct style) as nstyles  
FROM      Beers  
GROUP BY brewer  
HAVING    count(name) > 4  
ORDER BY brewer;
```

brewer	nbeers	nstyles
Bentspoke	9	7
Carlton	5	2
Frenchies	5	5
Hawkers	5	5
James Squire	5	4
One Drop	9	7
Sierra Nevada	5	5
Tallboy and Moose	5	5

distinct required, otherwise **nbeers=nstyles** for all brewers

❖ Filtering Groups (cont)

Alternative formulation of division using **GROUP-BY** and **HAVING**

Example: Find bars that each sell all of the beers Justin likes.

```
SELECT DISTINCT S.bar
FROM   Sells S, Likes L on (S.beer = L.beer)
WHERE  L.drinker = 'Justin'
GROUP BY S.bar
HAVING count(S.beer) =
        (SELECT count(beer) FROM Likes
         WHERE drinker = 'Justin');
```

❖ Partitions

Sometimes it is useful to

- partition a table into groups
- compute results that apply to each group
- use these results with individual tuples in the group

Comparison with **GROUP-BY**

- **GROUP-BY** produces one tuple for each group
- **PARTITION** augments each tuple with group-based value(s)
- can use other functions than aggregates (e.g. ranking)
- can use attributes other than the partitioning ones

❖ Partitions (cont)

Syntax for **PARTITION**:

```
SELECT attr1, attr2, ...,
       aggregate1 OVER (PARTITION BY attri),
       aggregate2 OVER (PARTITION BY attrj), ...
FROM   Table
WHERE  condition on attributes
```

Note: the *condition* cannot include the *aggregate* value(s)

❖ Partitions (cont)

Example: show each city with daily temperature and temperature range

Schema: *Weather(city,date,temperature)*

```
SELECT city, date, temperature
       min(temperature) OVER (PARTITION BY city) as lowest,
       max(temperature) OVER (PARTITION BY city) as highest
FROM   Weather;
```

Output: *Result(city, date, temperature, lowest, highest)*

❖ Partitions (cont)

Example showing **GROUP BY** and **PARTITION** difference:

```
SELECT city, min(temperature) max(temperature)
FROM Weather GROUP BY city
```

Result: one tuple for each city *Result(city,min,max)*

```
SELECT city, date, temperature as temp,
       min(temperature) OVER (PARTITION BY city),
       max(temperature) OVER (PARTITION BY city)
FROM Weather;
```

Result: one tuple for each temperature measurement.

Produced: 27 Feb 2021