

# Serializability

---

- Serializability
- Conflict Serializability
- Conflict Serializability Example
- View Serializability
- View Serializability Example

## ❖ Serializability

---

Serializable schedule:

- concurrent schedule for  $T_1..T_n$  with final state  $S$
- $S$  is also a final state of a possible serial schedule for  $T_1..T_n$

Abstracting this needs a notion of **schedule equivalence**.

Two common formulations of **serializability**:

- **conflict serializability** (read/write operations occur in the "right" order)
- **view serializability** (read operations see the correct version of data)

## ❖ Conflict Serializability

Consider two transactions  $T_1$  and  $T_2$  acting on data item  $X$ .

Possible orders for read/write operations by  $T_1$  and  $T_2$ :

$T_1$ first	$T_2$ first	Equiv?
$R_1(X) R_2(X)$	$R_2(X) R_1(X)$	yes
$R_1(X) W_2(X)$	$W_2(X) R_1(X)$	no
$W_1(X) R_2(X)$	$R_2(X) W_1(X)$	no
$W_1(X) W_2(X)$	$W_2(X) W_1(X)$	no

If  $T_1$  and  $T_2$  act on different data items, result is always equivalent.

## ❖ Conflict Serializability (cont)

---

Two transactions have a potential **conflict** if

- they perform operations on the same data item
- at least one of the operations is a write operation

In such cases, the order of operations affects the result.

If no conflict, can swap order without affecting the result.

If we can transform a schedule

- by swapping the order of non-conflicting operations
- such that the result is a serial schedule

then we say that the schedule is **conflict serializable**.

## ❖ Conflict Serializability (cont)

Example: transform a concurrent schedule to serial schedule

```

T1: R(A) W(A)      R(B)      W(B)
T2:      R(A)      W(A)      R(B) W(B)
swap
T1: R(A) W(A) R(B)      W(B)
T2:      R(A) W(A)      R(B) W(B)
swap
T1: R(A) W(A) R(B)      W(B)
T2:      R(A)      W(A) R(B) W(B)
swap
T1: R(A) W(A) R(B) W(B)
T2:      R(A) W(A) R(B) W(B)

```

## ❖ Conflict Serializability (cont)

---

Checking for conflict-serializability:

- show that ordering in concurrent schedule
- cannot be achieved in any serial schedule

Method for doing this:

- build a **precedence-graph**
- nodes represent transactions
- arcs represent order of action on shared data
- arc from  $T_1 \rightarrow T_2$  means  $T_1$  acts on  $X$  before  $T_2$
- a cycle indicates *not* conflict-serializable.

## ❖ Conflict Serializability Example


Example schedule which is not conflict serializable:

T1:	R(X)		R(Y)	W(X)		W(Y)
T2:		R(X)		W(X)		
T3:		R(X)				W(X)

attempted swaps

T1:		R(X)	W(X)		R(Y)	W(Y)
T2:		R(X)		W(X)		
T3:	R(X)				W(X)	

Precedence graph for the above schedule:



## ❖ View Serializability

---

View Serializability is

- an alternative formulation of serializability
- that is less conservative than conflict serializability (CS)  
(some safe schedules that are view serializable are not conflict serializable)

As with CS, it is based on a notion of schedule equivalence

- a schedule is "safe" if *view equivalent* to a serial schedule

The idea: if, across the two schedules ...

- they read the same version of a shared object
- they write the same final version of an object

then they are **view equivalent**



## ❖ View Serializability (cont)

Two schedules  $S$  and  $S'$  on  $T_1..T_n$  are **view equivalent** iff

- for each shared data item  $X$ 
  - if, in  $S$ ,  $T_j$  reads the initial value of  $X$ ,  
then, in  $S'$ ,  $T_j$  also reads the initial value of  $X$
  - if, in  $S$ ,  $T_j$  reads  $X$  written by  $T_k$ ,  
then, in  $S'$   $T_j$  also reads the value of  $X$  written by  $T_k$  in  $S'$
  - if, in  $S$ ,  $T_j$  performs the final write of  $X$ ,  
then, in  $S'$ ,  $T_j$  also performs the final write of  $X$

To check serializability of  $S$ ...

- find a serial schedule that is *view equivalent* to  $S$
- from among the  $n!$  possible serial schedules

## ❖ View Serializability Example

Example: consider the following concurrent schedule

```

T1:  R(A)  W(A)           R(B)           W(B)
T2:           R(A)           W(A)           R(B)  W(B)

```

If view serializable, the read/write behaviour must be like one of

1. T1: R(A) W(A) R(B) W(B)  
    T2: R(A) W(A) R(B) W(B)
2. T1: R(A) W(A) R(B) W(B)  
    T2: R(A) W(A) R(B) W(B)

## ❖ View Serializability Example (cont)

Reminder of concurrent schedule

T1:	R(A)	W(A)		R(B)		W(B)
T2:			R(A)		W(A)	R(B) W(B)

In the concurrent schedule

- A: T1 reads initial, T2 reads T1's write, T2 writes final
- B: T1 reads initial, T2 reads T1's write, T2 writes final

In T1;T2

- A: T1 reads initial, T2 reads T1's write, T2 writes final
- B: T1 reads initial, T2 reads T1's write, T2 writes final

So, concurrent schedule is view equivalent to T1;T2

Produced: 25 Mar 2021