


SQL Constraints, Updates and Queries

Consider the following data model for a business organisation and its employees:



Employees are uniquely identified by an id (eid), and other obvious information (name, age, ...) is recorded about each employee. An employee may work in several departments, with the percentage of time spent in each department being recorded in the **WorksIn** relation (as a number in the range 0.0-1.0, with 1.0 representing 100%). The percentages for a given employee may not sum to one if the employee only works part-time in the organisation. Departments are also uniquely identified by an id (did), along with other relevant information, including the id of the employee who manages the department.

Based on the ER design and the above considerations, here is a relational schema to represent this scenario:

```

create table Employees (
    eid    integer primary key,
    ename  text,
    age    integer,
    salary real,
    primary key (eid)
);
create table Departments (
    did    integer primary key,
    dname  text,
    budget real,
    manager integer references Employees(eid)
);
create table WorksIn (

```

```
eid      integer references Employees(eid),
did      integer references Departments(did),
percent  real,
primary key (eid,did)
);
```

Answer each of the following questions for this schema ...

1. Does the order of table declarations above matter?

Answer:

Yes. The order matters. We can not insert a Department tuple until there is an Employee tuple available to be the manager of the department. We cannot also insert any WorksIn tuple until you have both the Employee tuple and the Department tuple where the employee works.

2. A new government initiative to get more young people into work cuts the salary levels of all workers under 25 by 20%. Write an SQL statement to implement this policy change.

Answer:

SQL statement to reduce pay for people under 25 by 20%:

```
update Employees
set    salary = salary * 0.8
where  age < 25;
```

A straightforward application of the SQL UPDATE statement.

3. The company has several years of growth and high profits, and considers that the Sales department is primarily responsible for this. Write an SQL statement to give all employees in the Sales department a 10% pay rise.

Answer:

SQL statement to give Sales employees a 10% pay rise:

```
update Employees e
set    e.salary = e.salary * 1.10
where  eid in
        (select eid
         from   Departments d, WorksIn w
         where  d.dname = 'Sales' and d.did = w.did
        );
```

This query requires that we know which department an employee works for before updating their pay. The only information we have in the `Employees` relation to help with this is the employee id. Thus the subquery gives a list of ids for all employees working in the Sales department.

4. Add a constraint to the `CREATE TABLE` statements above to ensure that every department must have a manager.

Answer:

```
create table Departments (  
    did        integer,  
    dname      varchar(20),  
    budget     real,  
    manager    integer not null,  
    primary key (did),  
    foreign key (manager) references Employees(eid)  
);
```

Change the definition of `Departments` to ensure that the foreign key `manager` is not null. The foreign key constraint already ensures that the value must be a primary key in the `Employees` relation but, without the `NOT NULL` declaration, a foreign key value is allowed to be null.

5. Add a constraint to the `CREATE TABLE` statements above to ensure that no-one is paid less than the minimum wage of \$15,000.

Answer:

This is a straightforward example of a single-attribute constraint. It is effectively a constraint on the domain of the `salary` attribute. In fact, it would be sensible to have an initial constraint to ensure that the `salary` was at least positive. A similar constraint should probably be applied to the `percent` attribute in the `WorksIn` relation.

```
create table Employees (  
    eid        integer,  
    ename      varchar(30),  
    age        integer,  
    salary     real check (salary >= 15000),  
    primary key (eid)  
);
```

6. Add a constraint to the `CREATE TABLE` statements above to ensure that no employee can be committed for more than 100% of his/her time. Note that the SQL standard allows queries to be used in constraints, even though DBMSs don't implement this (for performance reasons).

Answer:

We have expressed this as a tuple-level constraint, even though it's really a constraint on the `eid` attribute. Note the use of scoping in the sub-query: the `eid` refers to the id of the employee that is being inserted or modified, while `w.eid` is bound by the tuple variable in the subquery. The subquery itself computes the total percent that the employee `eid` works.

This query is valid according to Ullman/Widom's description of the SQL2 standard. However, neither PostgreSQL nor Oracle supports subqueries in CHECK conditions. The condition can only be an expression involving the attributes in the updated row.

```
create table WorksIn (  
    eid            integer,  
    did            integer,  
    percent        real,  
    primary key (eid,did),  
    foreign key (eid) references Employees(eid),  
    foreign key (did) references Departments(did)  
    constraint MaxFullTimeCheck  
        check (1.00 >= (select sum(w.percent)  
                        from   WorksIn w  
                        where  w.eid = eid)  
                )  
);
```

In most relational database management systems, the constraint checking required here would need to be implemented via a trigger.

7. Add a constraint to the CREATE TABLE statements above to ensure that a manager works 100% of the time in the department that he/she manages. Note that the SQL standard allows queries to be used in constraints, even though DBMSs don't implement this (for performance reasons).

Answer:

We have expressed this as a tuple-level constraint, even though it's really a constraint on the `manager` attribute.

```
create table Departments (  
    did            integer,  
    dname          varchar(20),  
    budget         real,  
    manager        integer,  
    primary key (did),  
    foreign key (manager) references Employees(eid)  
    constraint FullTimeManager  
        check (1.0 = (select w.percent  
                      from   WorksIn w  
                      where  w.eid = manager)  
                )  
);
```

```
);
```

As in the previous question, this kind of constraint is allowed by the SQL standard, but DBMSs typically don't implement cross-table constraints like this; a trigger is required and the check has to be programmed in the trigger.

8. When an employee is removed from the database, it makes sense to also delete all of the records that show which departments he/she works for. Modify the CREATE TABLE statements above to ensure that this occurs.

Answer:

The ON DELETE CASCADE clause ensures that when the Employees record for `eid` is removed, then any WorksIn tuples that refer to `eid` are also removed.

```
create table WorksIn (
    eid          integer,
    did          integer,
    percent      real,
    primary key (eid,did),
    foreign key (eid) references Employees(eid) on delete cascade,
    foreign key (did) references Departments(did)
);
```

Of course, this immediately raises the issue of references to the Departments relation; this is considered in the next question.

9. When a manager leaves the company, there may be a period before a new manager is appointed for a department. Modify the CREATE TABLE statements above to allow for this.

Answer:

If the department has no manager, we indicate this by putting a value of NULL for the `manager` field. However, in one of the questions above, we added a NOT NULL constraint to ensure that every department *does* have a manager. To solve this question, we need to remove that constraint.

An alternative would be to always appoint a temporary manager, which could be accomplished via an UPDATE statement, e.g.

```
update department set manager = SomeEID where did = OurDeptID;
```

10. Consider the deletion of a department from a database based on this schema. What are the options for dealing with referential integrity between Departments and WorksIn? For each option, describe the required behaviour in SQL.

Answer:

Three possible approaches to referential integrity between Departments and WorksIn:

- a. Disallow the deletion of a Departments tuple if any Works tuple refers to it. This is the default behaviour, which would result from the CREATE TABLE definition in the previous question.
- b. When a Departments tuple is deleted, also delete all WorksIn tuples that refer to it. This requires adding an ON DELETE CASCADE clause to the definition of WorksIn.

```
create table WorksIn (
    eid      integer,
    did      integer,
    percent  real,
    primary key (eid,did),
    foreign key (eid) references Employees(eid) on delete cascade,
    foreign key (did) references Departments(did) on delete cascade
);
```

In this solution, we've added the same functionality to the eid field as well (see previous question).

- c. For every WorksIn tuple that refers to the deleted department, set the did field to the department id of some existing 'default' department. Unfortunately, Oracle doesn't appear to implement this functionality. If it did, the definition of WorksIn would change to:

```
create table WorksIn (
    eid      integer,
    did      integer default 1,
    percent  real,
    primary key (eid,did),
    foreign key (eid) references Employees(eid) on delete cascade,
    foreign key (did) references Departments(did) on delete set default
);
```

11. For each of the possible cases in the previous question, show how deletion of the Engineering department would affect the following database:

EID	ENAME	AGE	SALARY
1	John Smith	26	25000
2	Jane Doe	40	55000
3	Jack Jones	55	35000
4	Superman	35	90000
5	Jim James	20	20000

DID	DNAME	BUDGET	MANAGER
1	Sales	500000	2
2	Engineering	1000000	4

3 Service			200000	4
EID	DID	PCT_TIME		
1	2	1.00		
2	1	1.00		
3	1	0.50		
3	3	0.50		
4	2	0.50		
4	3	0.50		
5	2	0.75		

Answer:

- a. Disallow ... The database would not change. The DBMS would print an error message about referential integrity constraint violation.
- b. ON DELETE CASCADE ... All of the tuples in the WorksIn relation that have did = 2 are removed, giving:

DID	DNAME	BUDGET	MANAGER
1	Sales	500000	2
3	Service	200000	4

EID	DID	PCT_TIME
2	1	1.00
3	1	0.50
3	3	0.50
4	3	0.50

- c. ON DELETE SET NULL ... All of the tuples in the WorksIn relation that have did = 2 have that attribute modified to NULL, giving:

DID	DNAME	BUDGET	MANAGER
1	Sales	500000	2
3	Service	200000	4

EID	DID	PCT_TIME
1	NULL	1.00
2	1	1.00


3	1	0.50
3	3	0.50
4	NULL	0.50
4	3	0.50
5	NULL	0.75

d. ON DELETE SET DEFAULT ... All of the tuples in the WorksIn relation that have did = 2 have that attribute modified to the default department (1), giving:

DID	DNAME	BUDGET	MANAGER
1	Sales	500000	2
3	Service	200000	4

EID	DID	PCT_TIME
1	1	1.00
2	1	1.00
3	1	0.50
3	3	0.50
4	1	0.50
4	3	0.50
5	1	0.75

Consider the following data model for a a business that supplies various parts:



Based on the ER design and the above considerations, here is a relational schema to represent this scenario:

```

create table Suppliers (
    sid      integer primary key,
    sname    text,
    address  text
  )
  
```



```
);  
create table Parts (  
    pid      integer primary key,  
    pname    text,  
    colour   text  
);  
create table Catalog (  
    sid      integer references Suppliers(sid),  
    pid      integer references Parts(pid),  
    cost     real,  
    primary key (sid,pid)  
);
```

Write SQL statements to answer each of the following queries ...

Note1: all of these solutions have alternative formulations. If you think you have a better solution than the one(s) presented here, let me know.

Note2: a useful strategy, when developing an SQL solution to an information request, is to express intermediate results as views; this has been done in a few solutions here, but you might like to consider reformulating more of them with views, for clarity.

12. Find the *names* of suppliers who supply some red part.

Answer:

```
select S.sname  
from   Suppliers S, Parts P, Catalog C  
where  P.colour='red' and C.pid=P.pid and C.sid=S.sid
```

or

```
select sname  
from   Suppliers natural join Catalog natural join Parts  
where  P.colour='red'
```

13. Find the *sids* of suppliers who supply some red or green part.

Answer:

```
select C.sid  
from   Parts P, Catalog C  
where  (P.colour='red' or P.colour='green') and C.pid=P.pid
```

or

```
select sid
from   Catalog C natural join Parts P
where  (P.colour='red' or P.colour='green')
```

14. Find the *sids* of suppliers who supply some red part or whose address is 221 Packer Street.

Answer:

```
select S.sid
from   Suppliers S
where  S.address='221 Packer Street'
       or S.sid in (select C.sid
                    from   Parts P, Catalog C
                    where  P.color='red' and P.pid=C.pid
                    )
```

15. Find the *sids* of suppliers who supply some red part and some green part.

Answer:

```
select C.sid
from   Parts P, Catalog C
where  P.color='red' and P.pid=C.pid
       and exists (select P2.pid
                   from   Parts P2, Catalog C2
                   where  P2.color='green' and C2.sid=C.sid and P2.pid=C2.pid
                   )
```

or

```
(select C.sid
 from Parts P, Catalog C
 where P.color='red' and P.pid=C.pid
 )
intersect
(select C.sid
 from Parts P, Catalog C
 where P.color='green' and P.pid=C.pid
 )
```

16. Find the *sids* of suppliers who supply every part.

Answer:

```

select S.sid
from   Suppliers S
where  not exists((select P.pid from Parts P)
                  except
                  (select C.pid from Catalog C where C.sid=S.sid)
                )

```

or

```

select C.sid
from   Catalog C
where  not exists(select P.pid
                  from   Part P
                  where  not exists(select C1.sid
                                    from   Catalog C1
                                    where  C1.sid=C.sid and C1.pid=P.pid
                                  )
                )

```

17. Find the *sids* of suppliers who supply every red part.

Answer:

```

select S.sid
from   Suppliers S
where  not exists((select P.pid from Parts P where P.color='red')
                  except
                  (select C.pid from Catalog C where C.sid=S.sid)
                )

```

or

```

select C.sid
from   Catalog C
where  not exists(select P.pid
                  from   Part P
                  where  P.colour='red' and
                        not exists(select C1.sid
                                    from   Catalog C1
                                    where  C1.sid=C.sid and C1.pid=P.pid
                                  )
                )

```

18. Find the *sids* of suppliers who supply every red or green part.

Answer:

```

select S.sid
from   Suppliers S
where  not exists((select P.pid from Parts P
                    where P.color='red' or P.color='green')
                  except
                  (select C.pid from Catalog C where C.sid=S.sid)
                  )

```

or

```

select C.sid
from   Catalog C
where  not exists(select P.pid
                  from   Part P
                  where  (P.colour='red' or P.colour='green') and
                        not exists(select C1.sid
                                  from   Catalog C1
                                  where  C1.sid=C.sid and C1.pid=P.pid
                                  )
                  )

```

19. Find the *sids* of suppliers who supply every red part or supply every green part.

Answer:

```

(select S.sid
 from   Suppliers S
 where  not exists((select P.pid from Parts P where P.color='red')
                   except
                   (select C.pid from Catalog C where C.sid=S.sid)
                   )
 )
union
(select S.sid
 from   Suppliers S
 where  not exists((select P.pid from Parts P where P.color='green')
                   except
                   (select C.pid from Catalog C where C.sid=S.sid)
                   )
 )

```

20. Find pairs of *sids* such that the supplier with the first *sid* charges more for some part than the supplier with the second *sid*.

Answer:

```
select C1.sid, C2.sid
from   Catalog C1, Catalog C2
where  C1.pid = C2.pid and C1.sid != C2.sid and C1.cost > C2.cost
```

21. Find the *pids* of parts that are supplied by at least two different suppliers.

Answer:

```
select C.pid
from   Catalog C
where  exists(select C1.sid
                from   Catalog C1
                where  C1.pid = C.pid and C1.sid != C.sid
            )
```

22. Find the *pids* of the most expensive part(s) supplied by suppliers named "Yosemite Sham".

Answer:

```
select C.pid
from   Catalog C, Suppliers S
where  S.sname='Yosemite Sham' and C.sid=S.sid and
       C.cost >= all(select C2.cost
                     from   Catalog C2, Suppliers S2
                     where  S2.sname='Yosemite Sham' and C2.sid=S2.sid
                     )
```

23. Find the *pids* of parts supplied by every supplier at a price less than 200 dollars (if any supplier either does not supply the part or charges more than 200 dollars for it, the part should not be selected).

Answer:

```
select C.pid
from   Catalog C
where  C.price < 200.00
group by C.pid
having count(*) = (select count(*) from Suppliers);
```